

Image::Scale

High performance image resizing



Andy Grundman // August 18, 2015

David Maxim Micic - *ECO*

The Problem

- Image resizing libraries are designed for desktop-class systems
- Logitech Media Server needed to run on much slower systems such as these —>
- The Touch may also have to play perfect 24/96 audio at the same time the image cache is being built!



Netgear ReadyNAS Duo
240MHz Sparc clone, no FPU



Marvell SheevaPlug
1.2 GHz ARM9, no FPU



Squeezebox Touch
SD & USB storage support
Freescale i.MX35 533Mhz ARM11, slow FPU

Sparc?!

- Resizing a typical 1500x1500 JPEG album cover on this guy takes...

libgd	34s
ImageMagick	30s



To make matters worse....

- The server needs to cache thumbnails for every album in several different sizes:

40x40 64x64	On-device art
50x50 100x100	Web UI
75x75	iOS/Android



Arcane - Known/Learned

libgd & ImageMagick

Pros

High Quality

IM has over 14 different algorithms to choose from

Cons

Floating-point math

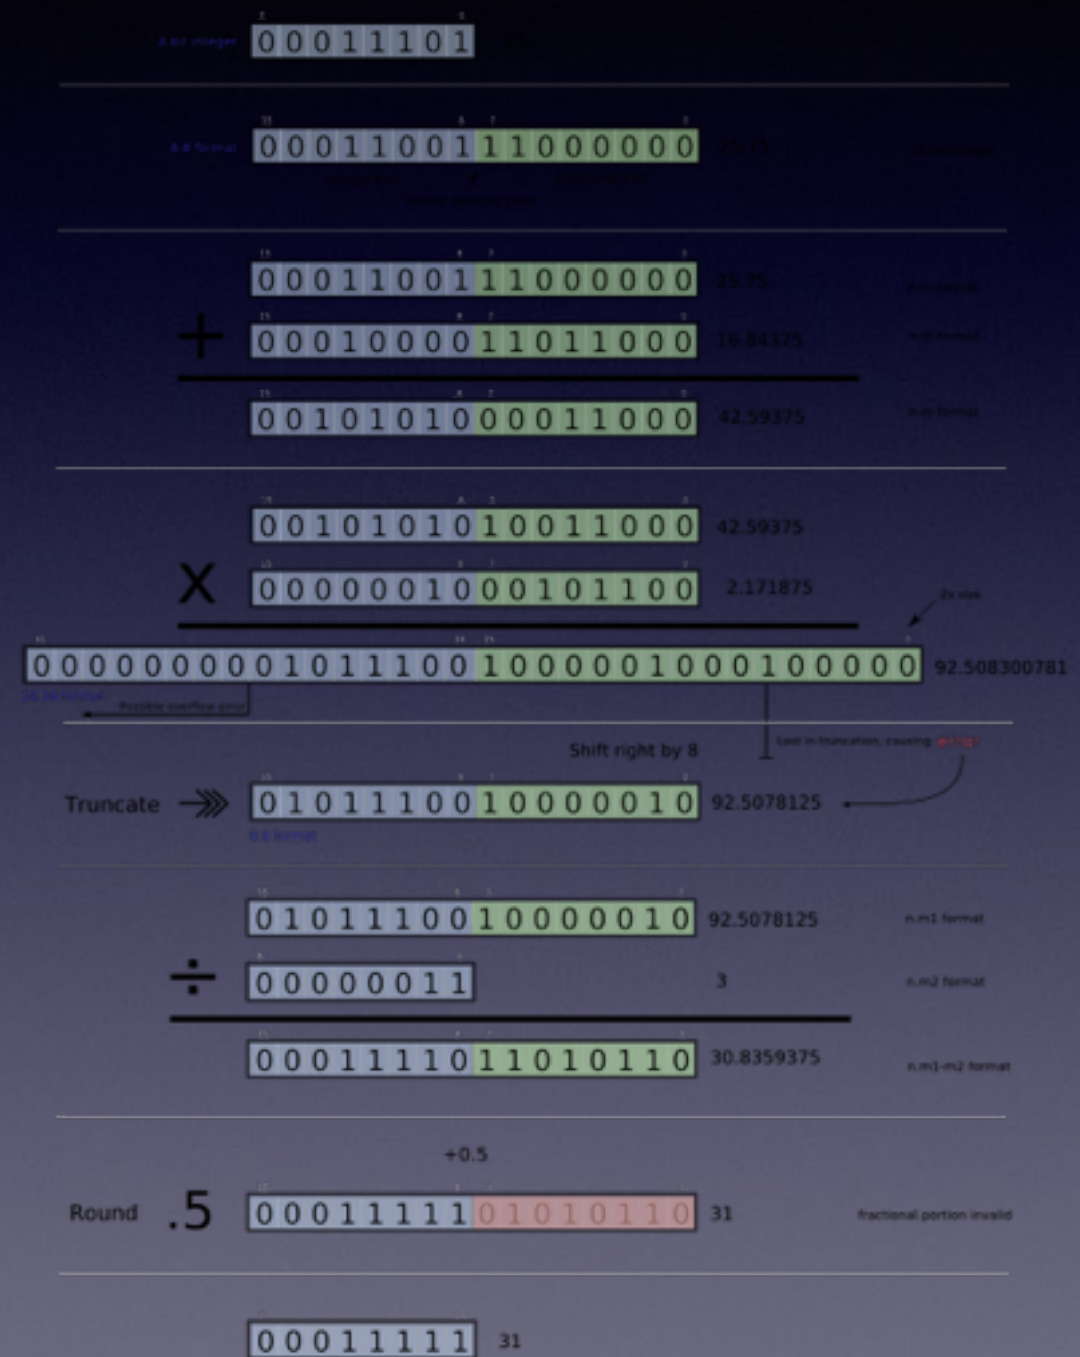
Inefficient memory use
(Touch LMS gets 64MB!)

No safeguards against loading massive images

Very difficult to build,
especially on Windows

The Solution

- Port the best quality resizing algorithms from both libgd and ImageMagick into fixed-point math.

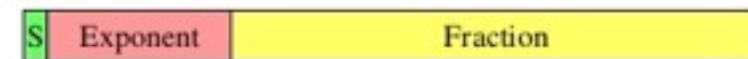


Fixed-Point Basics

- Stored in a 32-bit signed integer using 19.12 format
- First 19 bits are the integer part, and the remaining 12 are the fractional part. The remaining bit is the sign.
- Range of numbers that can be represented with 19.12:
0.000244140625 to 524287.999755859375
- Math operations retain ~4 decimal places of accuracy, resulting in essentially identical image quality.

Floating-Point Representation

- A floating-point number is represented by the triple
 - S is the Sign bit (0 is positive and 1 is negative)
 - Representation is called sign and magnitude
 - E is the Exponent field (signed)
 - Very large numbers have large positive exponents
 - Very small close-to-zero numbers have negative exponents
 - More bits in exponent field increases range of values
 - F is the Fraction field (fraction after binary point)
 - More bits in fraction field improves the precision of FP numbers



Value of a floating-point number = $(-1)^S \times \text{val}(F) \times 2^{\text{val}(E)}$

Converting to/from fixed

```
#define FRAC_BITS 12

typedef int32_t fixed_t;

static inline fixed_t int_to_fixed(int32_t x) {
    return x << FRAC_BITS;
}

static inline int32_t fixed_to_int(fixed_t x) {
    return x >> FRAC_BITS;
}

static inline fixed_t float_to_fixed(float x) {
    return ((fixed_t)((x) * (float)(1L << FRAC_BITS) + 0.5));
}

static inline float fixed_to_float(fixed_t x) {
    return ((float)((x) / (float)(1L << FRAC_BITS)));
}
```


Math Operations

```
// Note: gcc

#define FRAC_BITS 12

typedef int32_t fixed_t;

static inline fixed_t fixed_mul(fixed_t x, fixed_t y) {
    return (fixed_t)((((int64_t)x * y) >> FRAC_BITS));
}

static inline fixed_t fixed_div(fixed_t x, fixed_t y) {
    return (fixed_t)((((int64_t)x << FRAC_BITS) / y));
}
```

Mix in some ASM for even more performance!

```
// This improves fixed-point performance about 15-20% on x86
static inline fixed_t fixed_mul(fixed_t x, fixed_t y) {
    fixed_t __hi, __lo;
    __asm__ __volatile__(
        "imull %3\n"
        "shrdl %4, %1, %0"
        : "=a"(__lo), "=d"(__hi)
        : "%a"(x), "rm"(y), "I"(FRAC_BITS)
        : "cc"
    );
    return __lo;
}

// ARM
static inline fixed_t fixed_mul(fixed_t x, fixed_t y) {
    fixed_t __hi, __lo, __result;
    __asm__ __volatile__(
        "smull %0, %1, %3, %4\n\t"
        "movs %0, %0, lsr %5\n\t"
        "adc %2, %0, %1, lsl %6"
        : "=&r" (__lo), "=&r" (__hi), "=r" (__result)
        : "%r" (x), "r" (y), "M" (FRAC_BITS), "M" (32 - (FRAC_BITS))
        : "cc"
    );
    return __result;
}
```


Other cool features

- JPEG, PNG, and GIF support via libjpeg, libpng, and giflib. libjpeg-turbo (2-4x faster ASM-optimized version) is recommended.
- BMP is also supported, but I think only because I wanted to learn how BMP worked. :)
- Supports JPEG IDCT scaling, which efficiently returns a pre-shrunk version at certain fixed ratios such as 1/2, 1/4, 1/8, etc. Extremely important memory savings.
- Auto-rotates JPEG images that contain EXIF orientation metadata, e.g. iPhone photos.
- The floating-point versions of each resizer are included to make benchmarking easy.

Benchmarks (1/3)

GD copyResampled	1x
resize_gd	3.16x
resize_gd_fixed_point	3.1x

2.4 GHz MacBook Pro (2009 model)
1425x1425 JPEG -> 200x200
libjpeg-turbo v8 with scaling

- On a fast CPU, the original floating-point code edges out the fixed-point version.
- However, in this module it is still over 3x faster than libgd!

Benchmarks (2/3)

GD copyResampled	1x
resize_gd	2x
resize_gd_fixed_point	7.4x

Marvell SheevaPlug 1.2GHz ARM9
1425x1425 JPEG -> 200x200
libjpeg-turbo 6b with scaling

- On a system with no FPU, fixed-point shines

Benchmarks (3/3)

GD copyResampled	1x
resize_gd	1.1x
resize_gd_fixed_point	66x

240MHz Netgear ReadyNAS Duo (Sparc)
JPEG 1425x1425 -> 200x200
libjpeg 6b with scaling

- Not a typo

Image Quality

200x200, enlarged 2x



Float



Fixed

Where?

- <http://search.cpan.org/dist/Image-Scale/>
- <https://github.com/andygrundman/Image-Scale>

Questions?

